

SQL

Part1
Introduction

Relational data model

关系数据模型

基本概念与术语

Data models 数据模型

relational data model 关系数据模型

Schemas 模式


data independence 数据独立性

示例-课程管理系统

Consider building a course management system (CMS):

Entities (e.g., Students, Courses)

Relationships (e.g., Alice is enrolled in A course)





Students



File Edit View Insert Format Data Tools Add-ons Help All changes saved in Drive

↶ ↷ 🖨 📄 100% \$ % .0 ← .00 123 Arial 10 B I ⚡ A 🔍 ⌂

fx

| | A | B | C | D | E | F | G |
|----|---|----------------|--------------------|----------------|-------------------|---|---|
| 1 | | | | | | | |
| 2 | | | | | | | |
| 3 | | | | | | | |
| 4 | | | | | | | |
| 5 | | Student | SID | Address | | | |
| 6 | | Mickey | 40001 | 43 Toontown | | | |
| 7 | | Daffy | 40002 | 147 Main St | | | |
| 8 | | Donald | 50003 | 312 Escondido | | | |
| 9 | | Minnie | 50004 | 451 Gates | | | |
| 10 | | Pluto | 10008 | 97 Packard | | | |
| 11 | | | | | | | |
| 12 | | | | | | | |
| 13 | | | | | | | |
| 14 | | Course | Description | Room | Class size | | |
| 15 | | cs145 | Toon systems | Nvidia | 300 | | |
| 16 | | cs161 | Animation art | Gates 300 | 145 | | |
| 17 | | cs245 | Painting town rec | Packard 45 | 27 | | |
| 18 | | | | | | | |

‘Modeling’ the CMS

Logical Schema

Students (*sid: string, name: string, gpa: float*)

Courses (*cid: string, cname: string, credits: int*)

Enrolled (*sid: string, cid: string, grade: string*)

| sid | Name | Gpa |
|-----|------|-----|
| 101 | Bob | 3.2 |
| 123 | Mary | 3.8 |

Students

Corresponding
keys

| sid | cid | Grade |
|-----|-----|-------|
| 123 | 564 | A |

Enrolled

| cid | cname | credits |
|-----|-------|---------|
| 564 | 564-2 | 4 |
| 308 | 417 | 2 |

Courses

!Data model

Key concept

Relational model (aka tables)

Simple and most popular

Elegant algebra (E.F. Codd et al)

Every relation has a schema

Logical Schema: describes types, names

Physical Schema: describes data layout

Virtual Schema (Views): derived tables

Data model:

Organizing principle of data + operations

Schema :

Describes blueprint of table (s)

Data independence

Key concept

Logical Data Independence

Protection from changes in the Logical Structure
of the data

*i.e. Should not need to ask : Can we add a new
entity or attribute without rewriting the application*

Physical Data Independence

Protection from Physical Layout Changes

*i.e. Should not need to ask : Which disks are the
data stored on? Is the data indexed?*

One of the most important reasons to use a DBMS

SQL language

preview

Preview

SQL queries

[sqltutorial.org/sql-cheat-sheet](http://www.sqltutorial.org/sql-cheat-sheet)

QUERYING DATA FROM A TABLE

**SELECT c1, c2 FROM t;**

Query data in columns c1, c2 from a table

SELECT * FROM t;

Query all rows and columns from a table

SELECT c1, c2 FROM t**WHERE condition;**

Query data and filter rows with a condition

SELECT DISTINCT c1 FROM t**WHERE condition;**

Query distinct rows from a table

SELECT c1, c2 FROM t**ORDER BY c1 ASC [DESC];**

Sort the result set in ascending or descending order

SELECT c1, c2 FROM t**ORDER BY c1****LIMIT n OFFSET offset;**

Skip offset of rows and return the next n rows

SELECT c1, aggregate(c2)**FROM t****GROUP BY c1;**

Group rows using an aggregate function

SELECT c1, aggregate(c2)**FROM t****GROUP BY c1****HAVING condition;**

Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

SELECT c1, c2**FROM t1****INNER JOIN t2 ON condition;**

Inner join t1 and t2

SELECT c1, c2**FROM t1****LEFT JOIN t2 ON condition;**

Left join t1 and t2

SELECT c1, c2**FROM t1****RIGHT JOIN t2 ON condition;**

Right join t1 and t2

SELECT c1, c2**FROM t1****FULL OUTER JOIN t2 ON condition;**

Perform full outer join

SELECT c1, c2**FROM t1****CROSS JOIN t2;**

Produce a Cartesian product of rows in tables

SELECT c1, c2**FROM t1, t2;**

Another way to perform cross join

SELECT c1, c2**FROM t1 A****INNER JOIN t2 B ON condition;**

Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT c1, c2 FROM t1**UNION [ALL]****SELECT c1, c2 FROM t2;**

Combine rows from two queries

SELECT c1, c2 FROM t1**INTERSECT****SELECT c1, c2 FROM t2;**

Return the intersection of two queries

SELECT c1, c2 FROM t1**MINUS****SELECT c1, c2 FROM t2;**

Subtract a result set from another result set

SELECT c1, c2 FROM t1**WHERE c1 [NOT] LIKE pattern;**

Query rows using pattern matching %, _

SELECT c1, c2 FROM t**WHERE c1 [NOT] IN value_list;**

Query rows in a list

SELECT c1, c2 FROM t**WHERE c1 BETWEEN low AND high;**

Query rows between two values

SELECT c1, c2 FROM t**WHERE c1 IS [NOT] NULL;**

Check if values in a table is NULL or not

Table of Contents

1. SQL introduction & schema definitions
2. Basic single-table queries: SFW
3. Basic multiple-table queries: Joins

练习代码：SQL-1.ipynb

SQL Definitions

principles

What you will learn about in this section

1. What is SQL?
2. Basic schema definitions
3. Keys & constraints intro

SQL Introduction

- SQL is a standard language for querying and manipulating data
- SQL is a **very high-level** programming language
This works because it is optimized well!
- Many standards out there:
ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3), ...

SQL stands for
Structured
Query
Language

SQL is a...

- Data Manipulation Language (DML)

- Query one or more tables

- Insert/delete/modify tuples in tables

- Data Definition Language (DDL)

- Define relational schemata

- Create/alter/delete tables and their attributes

Set algebra

List: [1, 1, 2, 3]

Set: {1, 2, 3}

Multiset: {1, 1, 2, 3}

UNIONS

Set: {1, 2, 3} \cup { 2 } = { 1, 2, 3 }

Multiset: {1, 1, 2, 3} \cup { 2 } = { 1, 1, 2, 2, 3 }

Cross-product

{1, 1, 2, 3} * { y, z } =

{ <1, y>, <1, y>, <2, y>, <3, y>

<1, z>, <1, z>, <2, z>, <3, z>

}

A **multiset** is an unordered list (or: a set with multiple duplicate instances allowed)

i.e. no *next()*, etc.
methods!

Tables in SQL

Product

| PName | Price | Manuf |
|-------------|----------|------------|
| Gizmo | \$19.99 | GizmoWorks |
| Powergizmo | \$29.99 | GizmoWorks |
| SingleTouch | \$149.99 | Canon |
| MultiTouch | \$203.99 | Hitachi |

A relation or table is a multiset of tuples having the attributes specified by the schema

Let's break this definition down

Tables in SQL

Product

| PName | Price | Manuf |
|-------------|----------|------------|
| Gizmo | \$19.99 | GizmoWorks |
| Powergizmo | \$29.99 | GizmoWorks |
| SingleTouch | \$149.99 | Canon |
| MultiTouch | \$203.99 | Hitachi |

An attribute (or column) is a typed data entry present in each tuple in the relation

NB: Attributes must have an atomic type in standard SQL, i.e. not a list, set, etc.

Tables in SQL

Product

| PName | Price | Manuf |
|-------------|----------|------------|
| Gizmo | \$19.99 | GizmoWorks |
| Powergizmo | \$29.99 | GizmoWorks |
| SingleTouch | \$149.99 | Canon |
| MultiTouch | \$203.99 | Hitachi |

A tuple or row is a single entry in the table having the attributes specified by the schema

Also referred to sometimes as a Record

Tables in SQL

Product

| PName | Price | Manuf |
|-------------|----------|------------|
| Gizmo | \$19.99 | GizmoWorks |
| Powergizmo | \$29.99 | GizmoWorks |
| SingleTouch | \$149.99 | Canon |
| MultiTouch | \$203.99 | Hitachi |

The number of tuples is the cardinality of the relation

The number of attributes is the arity of the relation

Data Types in SQL

Atomic types:

Characters: CHAR(20), VARCHAR(50)

Numbers: INT, BIGINT, SMALLINT, FLOAT

Others: MONEY, DATETIME...

Every attribute must have an atomic type

Hence tables are flat

Table Schemas

The **schema** of a table is the table name, its attributes, and their types:

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)

A **key** is an attribute whose values are unique; we underline a key

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)

Key constraints

A **key** is a **minimal subset of attributes** that acts as a unique identifier for tuples in a relation

- A key is an implicit constraint on which tuples can be in the relation
- i.e. if two tuples agree on the values of the key, then they must be the same tuple!

Students(sid:string, name:string, gpa: float)

1. Which would you select as a key?
2. Is a key always guaranteed to exist?
3. Can we have more than one key?

Declaring Schema

Students(sid: *string*, name: *string*, gpa: *float*)

```
CREATE TABLE Students (
    sid CHAR(20) ,
    name VARCHAR(50) ,
    gpa float,
    PRIMARY KEY (sid),
)
```

NULL and NOT NULL

- To say “don’t know the value” we use **NULL**
NULL has (sometimes painful) semantics, more detail later

Students(sid:string, name:string, gpa: float)

| sid | name | gpa |
|-----|------|------|
| 123 | Bob | 3.9 |
| 143 | Jim | NULL |

Say, Jim just enrolled in his first class.

In SQL, we may constrain a column to be NOT NULL, e.g., “name” in this table

SQL查询操作依赖的数学结构是？

- A 连表List
- B 集合Set
- C 多集MultiSet
- D 数组Array

提交

2. Single - table queries 单表查询

What you will learn about in this section

The SFW(Select-From-Where expression) query

Other useful operators: LIKE, DISTINCT, ORDER BY

SQL Query

- Basic form (there are many many more bells and whistles)

```
SELECT <attributes>
```

```
FROM <one or more relations>
```

```
WHERE <conditions>
```

Call this a SFW query.

Simple SQL Query: Selection

Selection is the operation of filtering a relation's tuples on some condition

```
SELECT *
FROM Product
WHERE Category = 'Gadgets'
```

| PName | Price | Category | Manuf |
|-------------|----------|-------------|---------|
| Gizmo | \$19.99 | Gadgets | GWorks |
| Powergizmo | \$29.99 | Gadgets | GWorks |
| SingleTouch | \$149.99 | Photography | Canon |
| MultiTouch | \$203.99 | Household | Hitachi |



| PName | Price | Category | Manuf |
|------------|---------|----------|--------|
| Gizmo | \$19.99 | Gadgets | GWorks |
| Powergizmo | \$29.99 | Gadgets | GWorks |

Simple SQL Query: Projection

Projection is the operation of producing an output table with tuples that have a subset of their prior attributes

```
SELECT Pname, Price, Manufacturer  
FROM Product  
WHERE Category = 'Gadgets'
```

| PName | Price | Category | Manuf |
|-------------|----------|-------------|---------|
| Gizmo | \$19.99 | Gadgets | GWorks |
| Powergizmo | \$29.99 | Gadgets | GWorks |
| SingleTouch | \$149.99 | Photography | Canon |
| MultiTouch | \$203.99 | Household | Hitachi |



| PName | Price | Manuf |
|------------|---------|--------|
| Gizmo | \$19.99 | GWorks |
| Powergizmo | \$29.99 | GWorks |

Notation

Input Schema

Product(PName, Price, Category,
Manufacturer)

```
SELECT Pname, Price, Manufacturer  
FROM Product  
WHERE Category = 'Gadgets'
```



Output Schema

Answer(PName, Price, Manfacturer)

A Few Details

SQL **commands** are case insensitive:

Same: SELECT, Select, select

Same: Product, product

Values are **not**:

Different: ‘Seattle’, ‘seattle’

Use single quotes for constants:

‘abc’ - yes

“abc” - no

LIKE: Simple String Pattern Matching

```
SELECT *
FROM   Products
WHERE  PName LIKE '%gizmo%'
```

s **LIKE** p: pattern matching on strings

p may contain two special symbols:

- % = any sequence of characters
- _ = any single character

DISTINCT: Eliminating Duplicates

```
SELECT DISTINCT  
Category  
FROM Product
```



| Category |
|-------------|
| Gadgets |
| Photography |
| Household |

Versus

```
SELECT Category  
FROM Product
```



| Category |
|-------------|
| Gadgets |
| Gadgets |
| Photography |
| Household |

ORDER BY: Sorting the Results

```
SELECT      PName, Price, Manufacturer  
FROM        Product  
WHERE       Category='gizmo' AND Price > 50  
ORDER BY    Price, PName
```

Ordering is ascending,
unless you specify the
DESC keyword.

Ties are broken by the
second attribute on the
ORDER BY list, etc.

3. Multiple - table queries: JOIN 多表查询 : JOIN

What you will learn about in this section

JOINS

Inner JOINs

Outer JOINs

Joins

Product(PName, Price, Category,
Manufacturer)

Company(CName, StockPrice, Country)

Ex: Find all products under \$200 manufactured in Japan; return their names and prices.

Joins

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Several equivalent ways to write a basic join in SQL:

```
SELECT PName, Price  
FROM Product  
JOIN Company  
ON Manufacturer = Cname  
WHERE Price <= 200  
      AND Country='Japan'
```

```
SELECT PName, Price  
FROM Product, Company  
WHERE Manufacturer = CName  
      AND Country='Japan'  
      AND Price <= 200
```

A few more later on

Joins

Product

| PName | Price | Category | Manufacturer |
|-------------|-------|-------------|--------------|
| Gizmo | \$19 | Gadgets | GizmoWorks |
| Powergizmo | \$29 | Gadgets | GizmoWorks |
| SingleTouch | \$149 | Photography | Canon |
| MultiTouch | \$203 | Household | Hitachi |

Company

| CName | Stock Price | Country |
|------------|-------------|---------|
| GizmoWorks | 25 | USA |
| Canon | 65 | Japan |
| Hitachi | 15 | Japan |

```

SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
  
```

| PName | Price |
|-------------|-------|
| SingleTouch | \$149 |

Tuple Variable Ambiguity in Multi-Table

Person(name, address, worksfor)

Company(name, address)

1. SELECT DISTINCT name,
address
2. FROM Person, Company
3. WHERE worksfor = name

Which “address”
does this refer to?

Which name's??

Tuple Variable Ambiguity in Multi-Table

Both equivalent ways to resolve variable ambiguity

Person(name, address, worksfor)

Company(name, address)

SELECT DISTINCT Person.name,

Person.address

FROM Person, Company

WHERE Person.worksfor =

Company.name

SELECT DISTINCT p.name, p.address

FROM Person p, Company c

WHERE p.worksfor = c.name

Semantics of JOINS

```
SELECT x1.a1, x2.a2, ..., xn.ak
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn
WHERE Conditions(x1, ..., xn)
```

Note:

This is a *multiset* union

```
Answer = {}
for x1 in R1 do
    for x2 in R2 do
        ....
    for xn in Rn do
        if Conditions(x1, ..., xn)
        then Answer = Answer U {(x1.a1, x1.a2, ..., xn.ak)}
return Answer
```

Semantics of JOINS

- Take **cross product**

$$X = R \times S$$

- Apply **selections/conditions**

$$Y = \{(r, s) \text{ in } X \mid r.A == s.B\}$$

- Apply **projections** to get final output

$$Z = (y.A) \text{ for } y \text{ in } Y$$

```
SELECT R.A  
FROM R, S  
WHERE R.A = S.B
```

Recall: Cross product ($A \times B$) is the set of all unique tuples in A, B

Ex: $\{a,b,c\} \times \{1,2\}$


$$\begin{aligned} &= \{(a,1), (a,2), (b,1), (b,2), (c,1), \\ &\quad (c,2)\} \end{aligned}$$

= Filtering !

= Returning only *some* attributes

Remembering this order is critical to understanding the output of certain queries (see later on...)

An example of SQL semantics



RECAP: Joins

By default, joins in SQL are “**inner joins**”:

```
Product(name, category)  
Purchase(prodName, store)
```

1

```
SELECT Product.name, Purchase.store  
FROM Product  
JOIN Purchase ON Product.name = Purchase.prodName
```

2

```
SELECT Product.name, Purchase.store  
FROM Product, Purchase  
WHERE Product.name = Purchase.prodName
```



Both equivalent:
Both INNER JOINS!

INNER JOIN:

Product

| name | category |
|----------|----------|
| Gizmo | gadget |
| Camera | Photo |
| OneClick | Photo |

Purchase

| prodName | store |
|----------|-------|
| Gizmo | Wiz |
| Camera | Ritz |
| Camera | Wiz |

3

```
SELECT Product.name, Purchase.store  
FROM Product  
INNER JOIN Purchase  
ON Product.name = Purchase.prodName
```

Note: another equivalent way to
write an INNER JOIN!



| name | store |
|--------|-------|
| Gizmo | Wiz |
| Camera | Ritz |
| Camera | Wiz |

Outer Joins

- An **outer join** returns tuples from the joined relations that don't have a corresponding tuple in the other relations
 - I.e. If we join relations A and B on $a.X = b.X$, and there is an entry in A with $X=5$, but none in B with $X=5$...
 - A LEFT OUTER JOIN will return a tuple (a, NULL)!
- Left outer joins in SQL:

```
SELECT Product.name, Purchase.store  
FROM Product  
LEFT OUTER JOIN Purchase ON  
Product.name = Purchase.prodName
```

Now we'll get products even if they didn't sell

LEFT OUTER JOIN:

Product

| name | category |
|----------|----------|
| Gizmo | gadget |
| Camera | Photo |
| OneClick | Photo |

Purchase

| prodName | store |
|----------|-------|
| Gizmo | Wiz |
| Camera | Ritz |
| Camera | Wiz |

```
SELECT Product.name, Purchase.store  
FROM Product  
LEFT OUTER JOIN Purchase  
ON Product.name = Purchase.prodName
```



| name | store |
|----------|-------|
| Gizmo | Wiz |
| Camera | Ritz |
| Camera | Wiz |
| OneClick | NULL |

Other Outer Joins

- **Left outer join:**
 - Include the left tuple even if there's no match
- **Right outer join:**
 - Include the right tuple even if there's no match
- **Full outer join:**
 - Include the both left and right tuples even if there's no match

JOIN连接操作基于的数学运算是？

A

内积(Inner product)

B

交叉积(Cross Product)

提交

Outer Joins

- Left outer join:
 - Include the left tuple even if there's no match
- Right outer join:
 - Include the right tuple even if there's no match
- Full outer join:
 - Include the both left and right tuples even if there's no match

多表查询的（ JOIN ）连接操作有哪些？

- A Inner JOIN
- B Left JOIN
- C Right JOIN
- D Outer JOIN

提交

参考资料

- 斯坦福大学数据库课程
- CS145 : Data Management and Data Systems
- CS245 : Principles of Data-Intensive Systems

**THANK
YOU!**